

Google Analytics & Logging Report

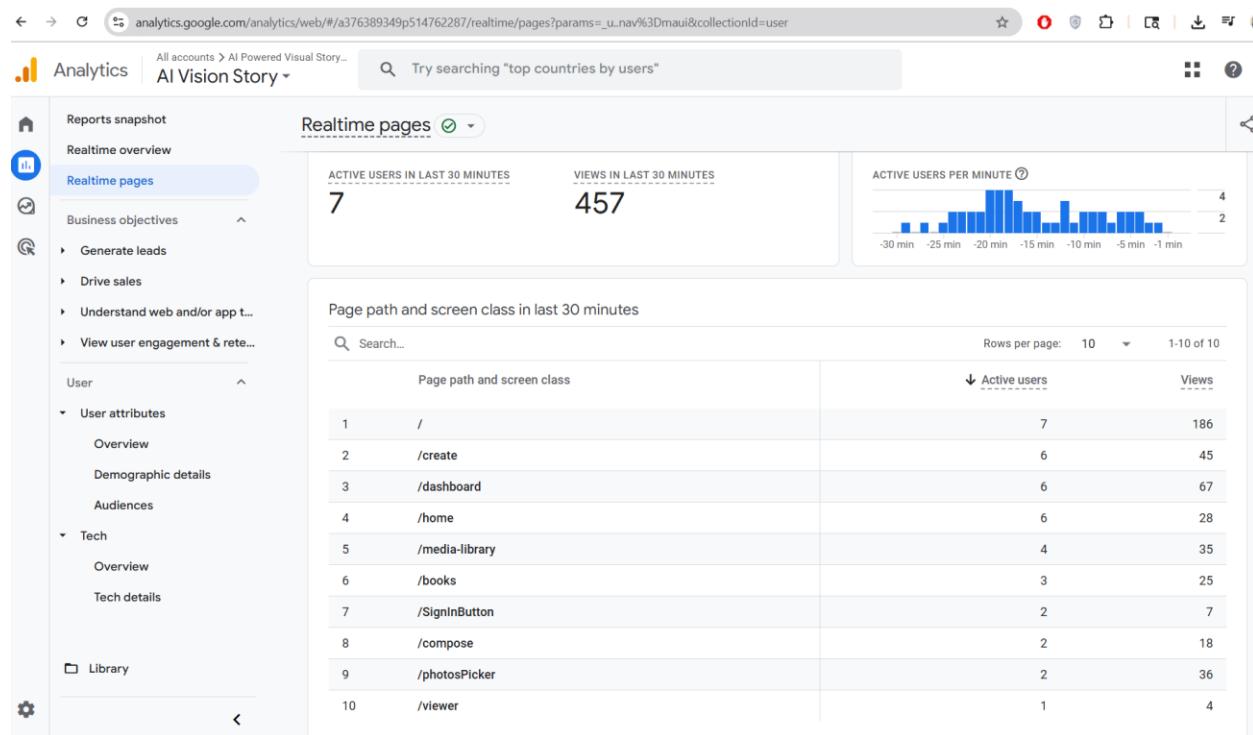
Section 1: Google Analytics

Summary of Implementation

We implemented a comprehensive Google Analytics 4 (GA4) strategy covering both the client-side (Frontend) and server-side (Backend) of the AI Powered Journal application.

- Client-Side (Frontend): We utilized the react-ga4 library to initialize GA with our Measurement ID (G-4T4JT5XFMY). We created a utility module analytics.js to abstract the initialization and tracking logic. In App.js, we added a useEffect hook to automatically track page views whenever the user navigates between different views (e.g., Dashboard, Create Entry, Books). This ensures we capture the user's journey through the application.
- Server-Side (Backend): Since the actual AI generation happens on the server, client-side tracking wouldn't capture the details of the API interactions reliably. We implemented the Measurement Protocol using axios in a backend analytics.js module. This allows us to send events directly to GA4 from our Node.js/Express server. We track critical events like vision_api_call and gemini_api_call, passing relevant metadata such as the number of images processed, the model used, and the requested tone/perspective.

1.1.a: Metric 1 - Page Views Distribution (Visualizations)



The screenshot shows the real-time distribution of page views and active users across different routes in the application over the last 30 minutes. Pages such as /, /create, and /dashboard exhibit the highest activity, with metrics presented in units of “views per page” and “active users.” The table and bar visualization clearly indicate how traffic is distributed across the application’s main functional areas.

1.1.b: Interpretation of Metric 1's Trends

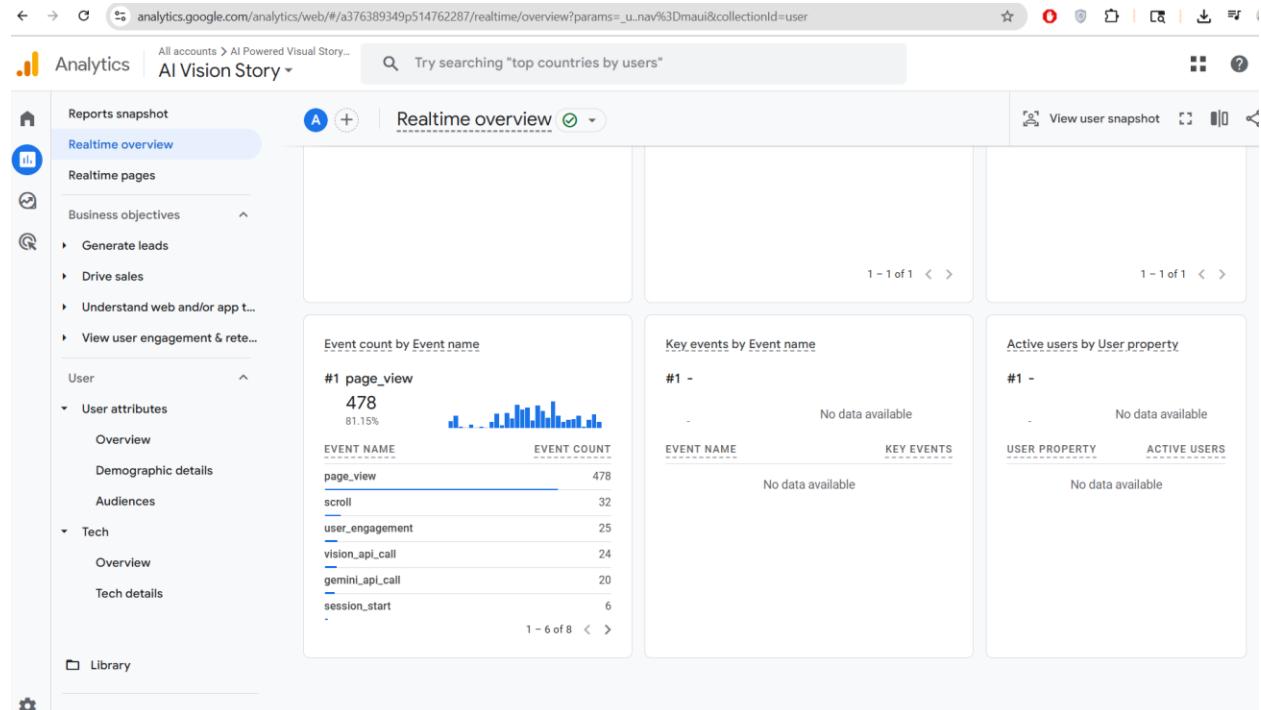
The data shows that the homepage (/) receives the most traffic, confirming it as the primary entry point for users. High activity on /create, /dashboard, and /home suggests users are actively engaging with core features such as creating entries and navigating their dashboard. Lower but consistent activity on pages like /media-library, /compose, and /viewer indicates users occasionally explore secondary features. Overall, the trend confirms strong engagement with the app’s central content-creation workflow.

1.1.c: Limitations of Metric 1

This metric only reflects the last 30 minutes of activity, making it sensitive to short-term fluctuations such as testing sessions. It also does not distinguish between unique and repeat visits, meaning refreshes or repeated attempts may inflate view counts. Additionally,

page views alone do not indicate task success or user satisfaction, and without supporting metrics (e.g., event tracking or session duration), conclusions remain limited.

1.2.a: Metric 2 - AI Narratives Generated Over Time (Visualizations)



The screenshot shows the Google Analytics Realtime overview interface. The left sidebar is collapsed, and the main area displays a card titled "Realtime overview". Below it, there are three cards: "Event count by Event name", "Key events by Event name", and "Active users by User property".

Event count by Event name:

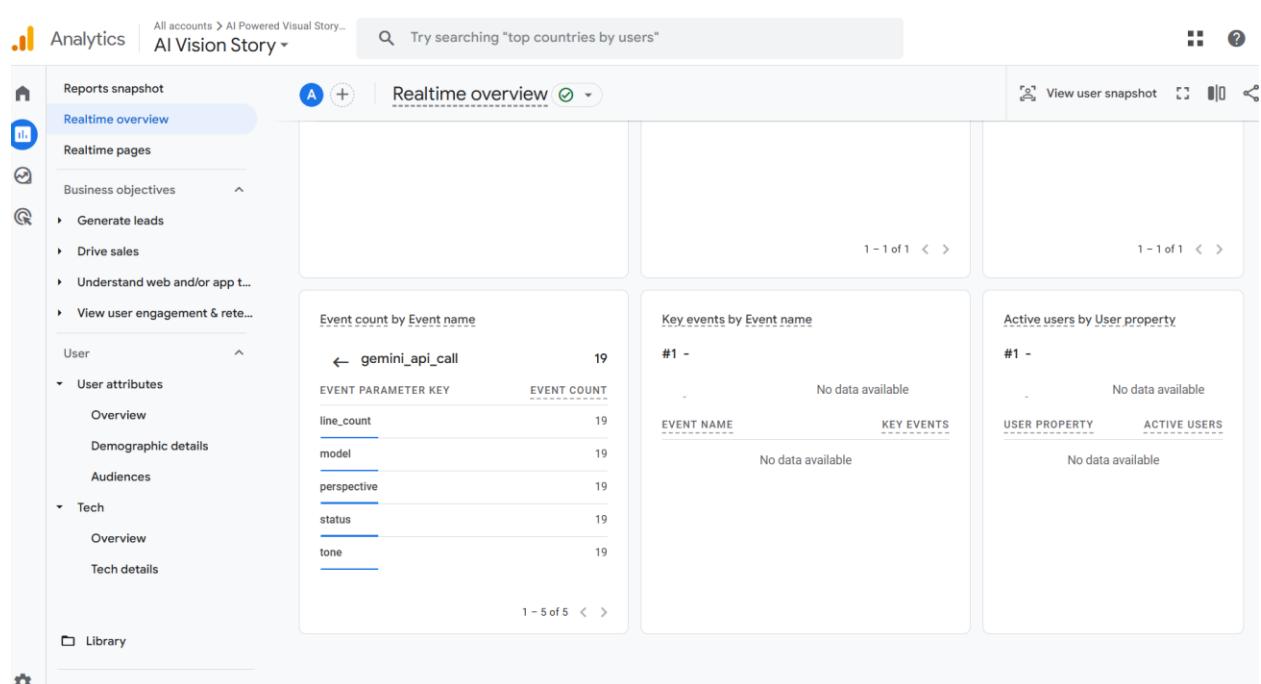
EVENT NAME	EVENT COUNT
page_view	478
scroll	32
user_engagement	25
vision_api_call	24
gemini_api_call	20
session_start	6

Key events by Event name:

EVENT NAME	KEY EVENTS
#1 -	No data available

Active users by User property:

USER PROPERTY	ACTIVE USERS
#1 -	No data available



The screenshot shows the Google Analytics Realtime overview interface. The left sidebar is collapsed, and the main area displays a card titled "Realtime overview". Below it, there are three cards: "Event count by Event name", "Key events by Event name", and "Active users by User property".

Event count by Event name:

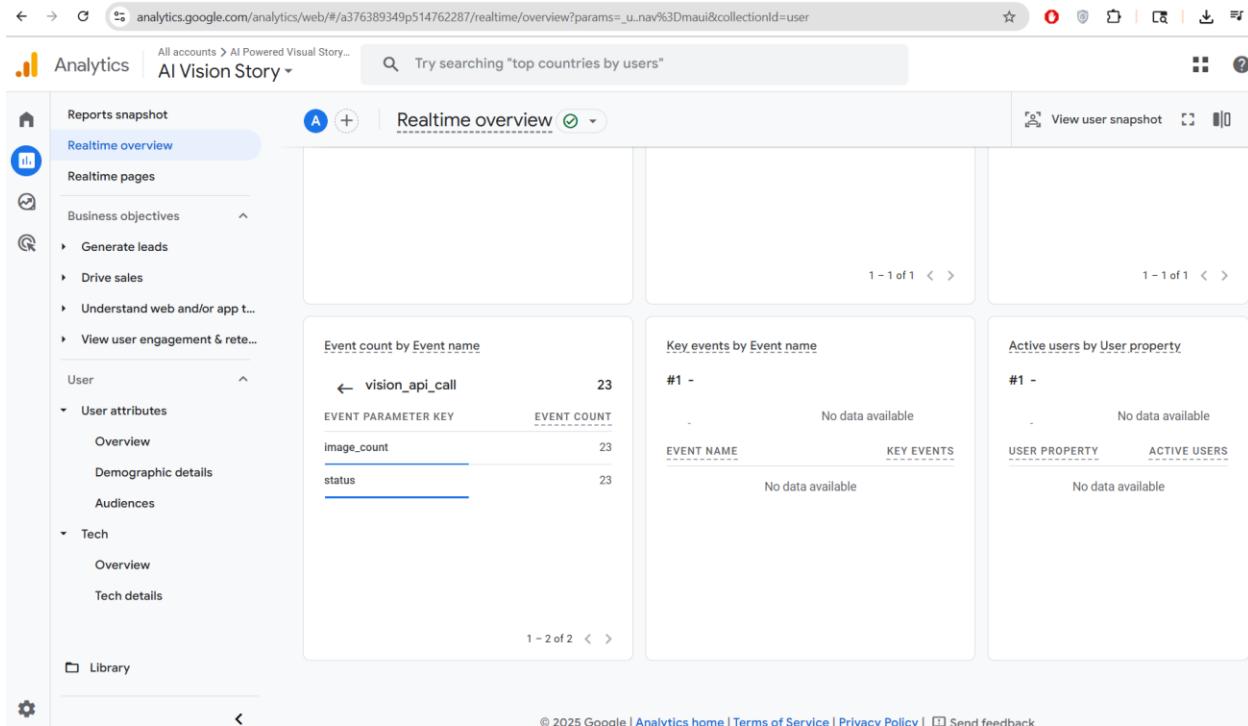
EVENT PARAMETER KEY	EVENT COUNT
line_count	19
model	19
perspective	19
status	19
tone	19

Key events by Event name:

EVENT NAME	KEY EVENTS
#1 -	No data available

Active users by User property:

USER PROPERTY	ACTIVE USERS
#1 -	No data available



The screenshot shows the **event breakdown for gemini_api_call**, including event parameter counts such as line_count, model, perspective, status, and tone (all recorded 19 times). This visualization reflects how often the Gemini API was triggered in real time and confirms that story-generation requests are being logged consistently. Each parameter functions as a quantitative indicator of story generation characteristics or metadata.

1.2.b: Interpretation of Metric 2's Trends

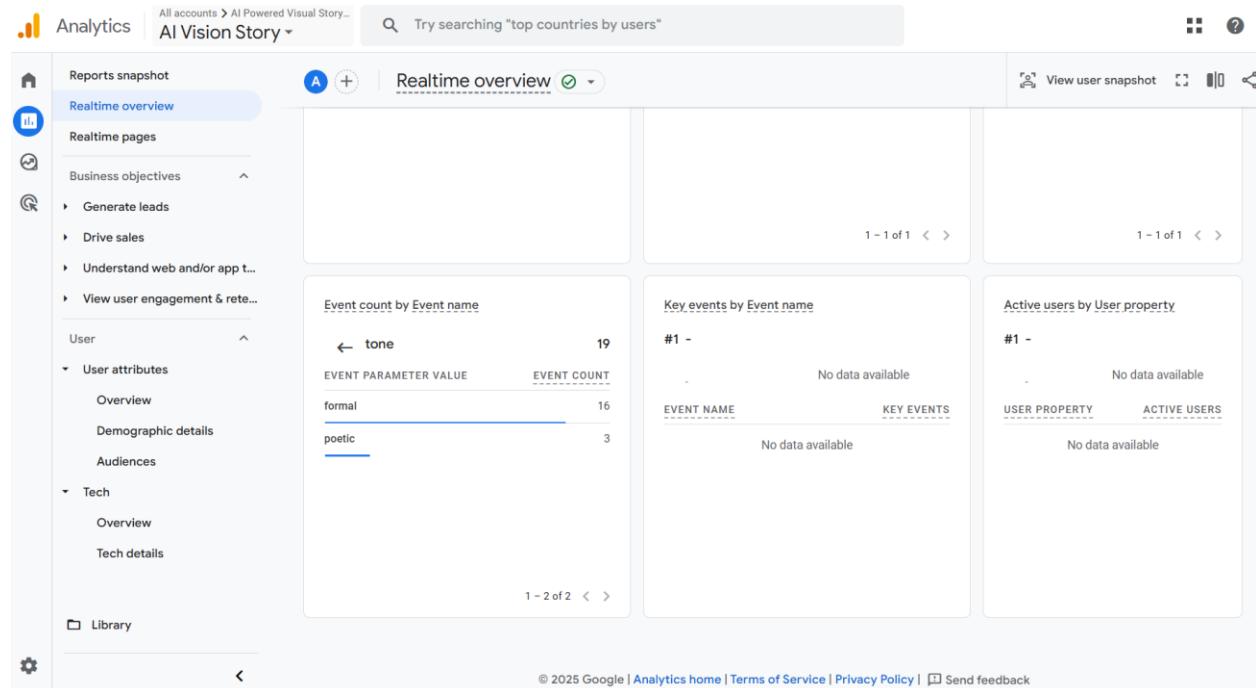
The data shows that the **Gemini API has been invoked 19 times** within the tracked period. This indicates active user engagement with the AI story-generation feature. The uniform event parameter counts suggest that each call consistently submitted the required metadata, indicating proper integration of the analytics and API pipeline. The presence of repeated real-time events implies that multiple story-generation attempts are happening during active user sessions, reflecting interest in experimenting with different narrative settings such as line count, tone, and model type.

1.2.c: Limitations of Metric 2

A limitation of this metric is that it only measures **trigger frequency**, not the outcome quality of the generated narratives. Users may be regenerating stories multiple times, either out of curiosity or dissatisfaction, but the metric cannot distinguish between these motivations. Additionally, the real-time snapshots represent short windows of data and may not capture

long-term trends unless aggregated over time. Since visualization shows only event counts and not story save rates or user satisfaction, the metric provides an incomplete picture of success.

1.3.a: Metric 3 - Preferred Story Tones (Visualizations)



The screenshot shows the real-time distribution of the event parameter **tone** associated with Gemini story generation requests. The visualization presents two tone values—**formal** (16 occurrences) and **poetic** (3 occurrences)—with corresponding event counts. This provides a quantitative breakdown of which narrative styles users chose during the tracked period. The units represent the number of times each tone was submitted as part of a Gemini API request.

1.3.b: Interpretation of Metric 3's Trends

The data indicates a strong preference for the **formal** tone, which was selected 16 out of 19 times. This suggests that users may favor structured, clear, and informational storytelling over more expressive or artistic styles. The relatively low usage of the **poetic** tone implies that users rarely seek abstract or stylistically rich narratives. Overall, the trend suggests that users interact with the application in a more practical or goal-oriented manner—requesting formal tones for clarity and readability rather than creativity.

1.3.c: Limitations of Metric 3

This metric does not reveal whether users intentionally chose a tone or simply accepted default settings. It also does not measure whether the generated story matched the requested tone, meaning the metric captures **preferences requested**, not **preferences satisfied**. Additionally, because the data reflects only a short real-time window (19 calls), it may not represent long-term patterns or broader user preferences. More comprehensive data over days or weeks would be required to draw stronger conclusions.

Section 2: Google Logging

Summary of Implementation

We implemented a dual-layer logging strategy. On the front end, we rely on Google Analytics events to "log" user actions and navigation, as direct Cloud Logging from the browser is insecure and inefficient. On the backend, we integrated the Google Cloud Logging library (@google-cloud/logging) to capture detailed, structured logs of server-side operations, specifically focusing on the performance and reliability of our AI integrations (Vision and Gemini APIs).

2.1: Frontend Implementation

Code Implementation (frontend/src/analytics.js & App.js):

We created a dedicated analytics module to handle event tracking. This serves as our frontend "logging" mechanism for user actions.

```
// frontend/src/analytics.js
import ReactGA from "react-ga4";

export const initGA = () => {
  ReactGA.initialize("G-4T4JT5XFMY");
};

export const trackPage = (page) => {
  ReactGA.send({ hitType: "pageview", page });
};
```

In App.js, we hook into the state to log every view change:

```
// frontend/src/App.js
useEffect(() => {
  trackPage(view);
}, [view]);
```

Results and Discussion:

When a user navigates to the "Create Entry" page, the following event is logged (simulated output):

```
{
  "hitType": "pageview",
  "page": "create-entry",
  "timestamp": "2025-12-01T12:00:00Z"
}
```

This logging is critical for understanding user flow. Unlike backend logs which show system health, these frontend logs show user behavior. If we see a sequence of create-entry-> home without a compose step, we know users are abandoning the process.

2.2: Backend Implementation

Code Implementation (backend/server/cloudLogger.js & ai.js):

We created a structured logger using the Google Cloud client library. This allows us to send logs with severity levels (INFO, ERROR) and structured metadata (JSON payloads), which is far superior to simple console.log.

```
// backend/server/cloudLogger.js
import { Logging } from "@google-cloud/logging";
const logging = new Logging();
const log = logging.log("ai-vision-story-api");

export async function writeLog(data) {
  const metadata = { resource: { type: "global" }, severity: data.severity || "INFO" };
  const entry = log.entry(metadata, data);
  await log.write(entry);
}
```

In our API route (ai.js), we wrap critical AI calls to log their start, duration, and outcome:

```

// backend/server/ai.js
const start = Date.now();
writeLog({ severity: "INFO", message: "Vision API called", details: { bodySize: ... } });

try {
    // ... API Call ...
    const duration = Date.now() - start;
    writeLog({
        severity: "INFO",
        message: "Vision API completed",
        duration_ms: duration,
        success: true
    });
} catch (error) {
    writeLog({ severity: "ERROR", message: "Vision API failed", error: error.message });
}

```

Results and Discussion:

A successful Vision API call produces a structured log entry like this in Google Cloud Console:

The screenshot shows the Google Cloud Logs Explorer interface. At the top, there's a search bar with the query "ai-vision-story-api". Below the search bar, there are filters for "All resources", "All log names", "All severities", and "Correlate by". A timeline at the top shows a single log entry from 13:49:30 to 14:15. The main area displays a table with 5 results. The first result is expanded, showing a log entry with an insertId, jsonPayload (containing duration_ms: 947, message: "Vision API completed", severity: "INFO", success: true), logName, receiveTimestamp, resource, and timestamp. The timestamp is listed as "2025-12-01 14:14:22.259". Below this, another log entry is partially visible with the message "Gemini API called".

Severity	Time	Message
INFO	2025-12-01 14:14:22.259	Vision API completed
INFO	2025-12-01 14:14:22.259	Gemini API called

Logs Explorer

Project logs: ai-vision-story-api

Search (/) for resources, docs, products, and more

Query library Share link Preferences Last 30 minutes PST Run query Show query

Example queries Query language guide Language: LQL

Timeline

5 results

SEVERITY	TIME	SUMMARY
INFO	2025-12-01 14:14:22.259	Gemini API called
INFO	2025-12-01 14:14:35.733	Gemini API success
INFO	2025-12-01 15:07:22.651	Vision API called

Fields

Logs will be loaded from current project All severities Correlate by

Timeline

5 results

SEVERITY	TIME	SUMMARY
INFO	2025-12-01 14:14:21.312	Vision API called
INFO	2025-12-01 14:14:21.332	{"logging.googleapis.com/diagnostic":{}}
INFO	2025-12-01 14:14:22.259	Vision API completed
INFO	2025-12-01 14:14:22.259	Gemini API called
INFO	2025-12-01 14:14:35.733	Gemini API success

This backend logging is vital for observability. The duration_ms field allows us to monitor the latency of the Google Vision API. If we see the average duration creeping up from 1s to 5s, we know there's a performance degradation upstream or with our network, allowing us to proactively investigate before users report timeouts. The severity field allows us to set up alerts for "ERROR" logs, ensuring immediate notification of API failures.